

Predicting Future Euro Exchange Rates Using a Random Forest Regressor

Objective

Goal: To predict future Euro exchange rates using historical data and a Random Forest Regressor.

Deliverables:

1. EDA and Preprocessing: Show the handling of missing data, feature selection, and creation of lagged variables.
2. Model Training and Evaluation: Present the results of training the Random Forest Regressor and provide evaluation metrics (MAE, MSE, R^2).
3. Visualization: Plot of actual vs predicted Euro exchange rates.

Data and Tools Used

Dataset: Historical Euro exchange rate data.

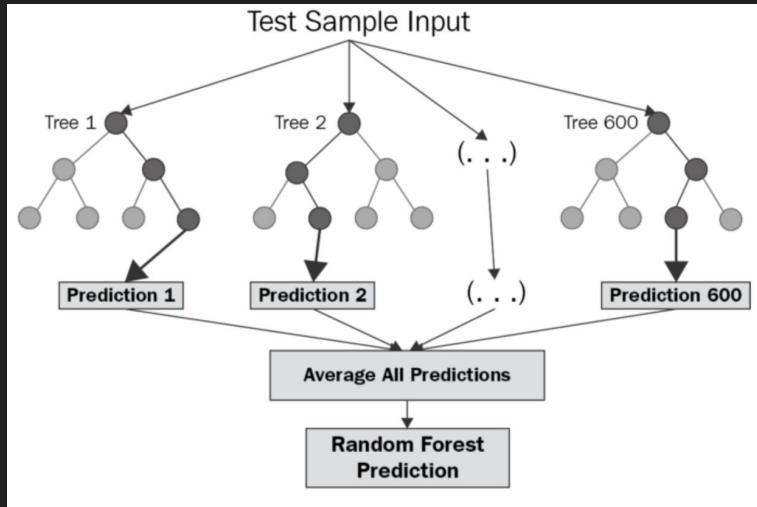
Tools:

- Python: Language used.
- Libraries: Pandas, NumPy, Scikit-Learn, Matplotlib
- Data Points: Daily exchange rates, used to predict the next day's rate.

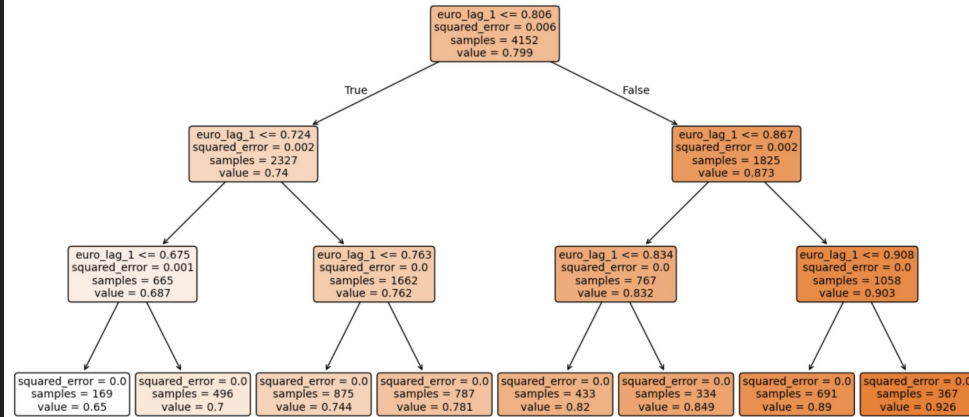
What is Random Forest Regressor?

- Ensemble Learning: Combines multiple decision trees to improve accuracy and reduce errors.
- Decision Trees: Each tree in the forest learns from different subsets of the data.
- Averaging Predictions: By averaging outputs from all trees, Random Forest reduces variance and enhances stability.
- Suitable for predicting trends in time series data, like exchange rates.

What is Random Forest Regressor?



Simplified Decision Tree for Euro Exchange Rate Prediction (Depth = 3)



<https://www.keboola.com/blog/random-forest-regression>

Setting Up the Environment and Loading Data

Code Overview: This code sets up the environment, mounts Google Drive, and loads the dataset for analysis.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.inspection import PartialDependenceDisplay
```

```
from google.colab import drive
drive.mount('/content/drive', force_remount=True)
```

Mounted at /content/drive

```
[3] data = pd.read_csv('/content/drive/My Drive/exchange_rates.csv')

print("Initial Data Overview:")
print(data.head())
print(data.info())
```

Data columns (total 52 columns):

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	5220 non-null	object
1	chinese_yuan	4845 non-null	float64
2	euro	5152 non-null	float64
3	japanese_yen	4935 non-null	float64
4	uk_pound	5126 non-null	float64
5	us_dollar	5220 non-null	float64
6	algerian_dinar	3328 non-null	float64
7	australian_dollar	4989 non-null	float64

EDA and Preprocessing

- Handling Missing Data: Used forward fill to handle any gaps in the dataset.

```
# Handle missing values with forward fill
data.fillna(method='ffill', inplace=True)
print("\nMissing values after forward-fill:")
print(data.isnull().sum())
```

- Feature Selection: Focused on the euro column to build our target variable.
- Creation of Lagged Variable:
 - Created euro_lag_1, a 1-day lagged value, as the main feature to capture short-term trends.

```
euro_data['euro_lag_1'] = euro_data['euro'].shift(1)
```

	B	C	D	E	F	G	
1	chinese_yuan	euro	japanese_yen	uk_pound	us_dollar	algerian_dinar	a
2	32	0.7941550190597200		0.5599731212901780	1.0		1
3	35	8.277 0.790076637433831	106.9	0.5556790397866190	1.0		1
4	36	8.2771 0.7839448102853560	106.25	0.5491778977428744	1.0		1
5	37	8.2772 0.7887057338906850	106.24	0.5512679162072770	1.0		1
6	38	8.2772 0.791514959632737	106.15	0.5518154729058600	1.0		1
7	39	8.277 0.7851142341210650	107.55	0.5445140212360470	1.0		1
8	42	8.277 0.7795447458684130		0.5398693516169090	1.0		1
9	43	8.2768 0.7843752451172640	106.53	0.5420054200542010	1.0		1
10	44	8.2768 0.7878978884336590	106.22	0.5446623093681920	1.0		1
11	45	8.2767 0.7914523149980210	106.25	0.5483658697082690	1.0		1
12	46	8.2767 0.8004482510205720	106.1	0.5520287054926960	1.0		1
13	49	8.2768 0.8082114281095940	106.55	0.5599417660563300	1.0		1
14	50	8.277 0.797702616464582	107.5	0.5531585352361990	1.0		1
15	51	8.2771 0.7932101213611490	107.18	0.5463286713286710	1.0		1
16	52	8.2771 0.7867820613690010	106.88	0.5423581733376720	1.0		1
17	53	8.2771 0.7878978884336590	106.2	0.5412426932236420	1.0		1
18	56	8.2771 0.7952286282306160	106.22	0.547075879444760	1.0		1
19	57	8.2771 0.7989134776703880	106.15	0.5535259603675410	1.0		1
20	58	8.2771 0.7959882193743530	106.2	0.5452562704471100	1.0		1
21	59	8.277 0.8020532563362210	106.1	0.5480055080505040	1.0		1
22	60	8.2768 0.8074935400516800	105.97	0.552242102937928	1.0		1
23	62	8.277 0.8025038118931070	105.61	0.54878717803314680	1.0		1
24	63	8.2772 0.7945967421533570	105.53	0.544069640914037	1.0		1
25	64	8.2772 0.7984669434685400	105.39	0.5441288497116120	1.0		1

Model Training

Model Chosen: Random Forest Regressor

Reason: Random Forest is effective for time series prediction as it leverages multiple decision trees to capture trends.

Training Process:

- Split data into training (80%) and testing (20%) sets.
- Trained the model using only the 1-day lag as the primary feature.

```
0s # Define the feature (only 'euro_lag_1') and target
X = euro_data[['euro_lag_1']] # Only the 1-day lag
y = euro_data['euro']

# Split data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=False)

0s # Initialize and train the Random Forest Regressor
rf_model = RandomForestRegressor(n_estimators=100, random_state=42, oob_score=True)
rf_model.fit(X_train, y_train)
```

RandomForestRegressor

RandomForestRegressor(oob_score=True, random_state=42)

Model Evaluation Metrics

Out-of-Bag (OOB) Score: Used to assess the model's generalization.

Mean Absolute Error (MAE): Indicates average error in prediction.

Mean Squared Error (MSE): Measures the variance of prediction errors.

R² Score: Measures the percentage of variance explained by the model.

- Higher R² suggests better model performance.

```
▶ y_pred = rf_model.predict(X_test)


oob_score = rf_model.oob_score_
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f'Out-of-Bag Score: {oob_score}')
print("Mean Squared Error (MSE):", mse)
print('Mean Absolute Error (MAE):', mae)
print(f'R-squared Score: {r2}')
```

```
↔ Out-of-Bag Score: 0.9948508067815014
Mean Squared Error (MSE): 0.0001995978640355459
Mean Absolute Error (MAE): 0.00704736401456525
R-squared Score: 0.9207432390481782
```

Visualization of Results

Actual vs. Predicted Plot:



Key Insights and Conclusion

Key Takeaways:

- The model captured short-term patterns well, leveraging recent values.
- High R^2 and low errors indicate accurate predictions for this task.

Conclusion:

- Random Forest is effective for predicting exchange rates with minimal feature engineering.
- This model could be enhanced with more features or by using an ensemble approach for better accuracy.

Resources:

<https://www.keboola.com/blog/random-forest-regression>

<https://scikit-learn.org/dev/modules/generated/sklearn.ensemble.RandomForestRegressor.html>